

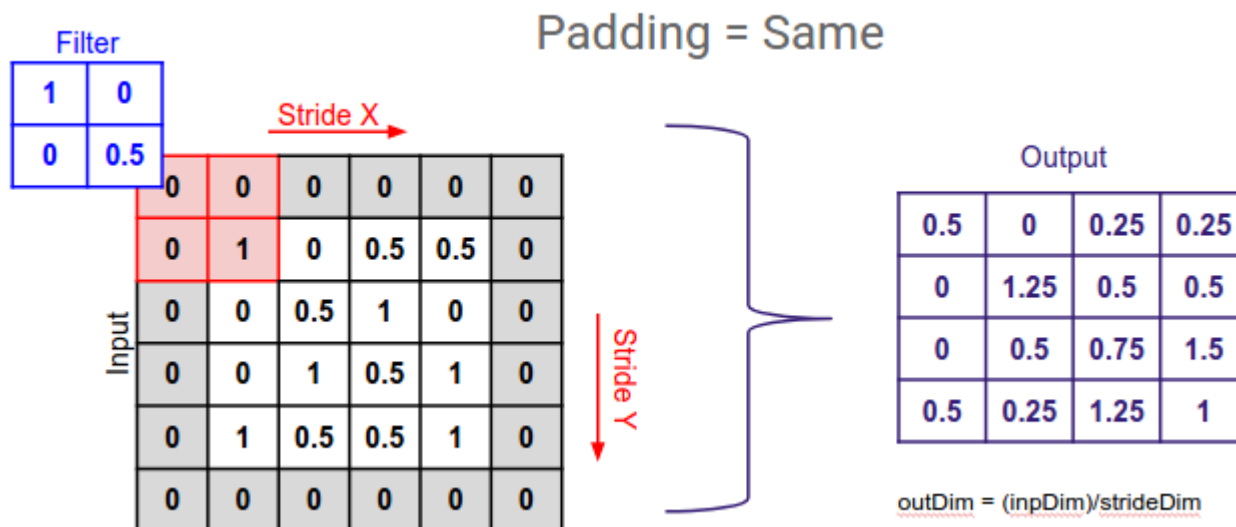


AlexNet

Computer Vision & Augmented Reality 연구실  
학부연구생 강 준 구

# Why do we use the padding in CNN?

- ▶ zero-padding
  - ▶ It conserves the pixel information on the edge side
  - ▶ It can preserve the input's spatial size



# AlexNet(2012)

## ► ILSVRC



# Contents

---

- ▶ 1. DataSet
- ▶ 2. Activation Function
- ▶ 3. Training on Multiple GPUs
- ▶ 4. Reducing Overfitting
- ▶ 5. AlexNet Architecture
- ▶ 6. Future work

# DataSet



- ▶ An RGB image of size 256 x 256
  - ▶ If the input image is not 256 x 256 or 3-channel RGB
    - ▶ It needs to be converted to 256 x 256 before using it for training the network
    - ▶ It needs to be converted to an RGB image



# Activation Function


## ► ReLU

- ReLUs train several times faster than their equivalents with tanh units.



# Activation function

Vanishing gradient (NN winter2: 1986-2006)



Activation functions on CIFAR-10

maxout	ReLU	VReLU	tanh	Sigmoid
<b>93.94</b>	<b>92.11</b>	92.97	89.28	n/c
93.78	91.74	92.40	89.48	n/c
–	91.93	<b>93.09</b>	–	n/c
91.75	90.63	92.27	<b>89.82</b>	n/c
n/c†	90.91	92.43	89.54	n/c


# Training on Multiple GPUs

## ► GTX 580 3GB GPU

► 5~6 days

► 90 epochs

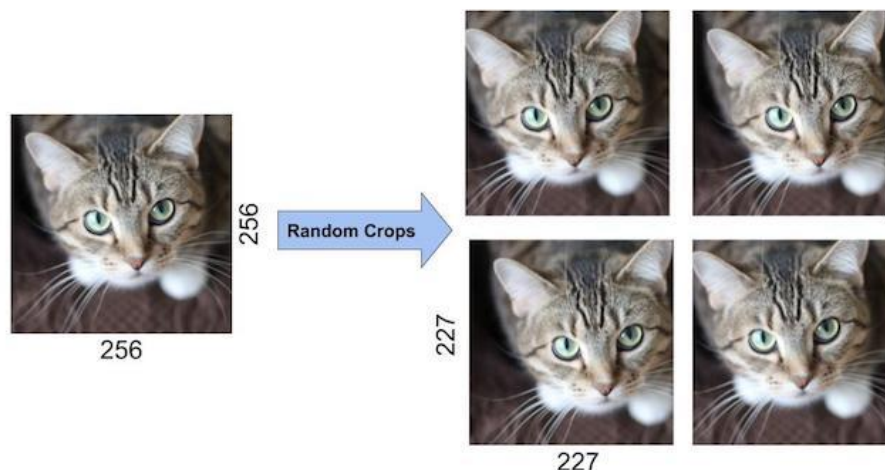
□ The results can be improved





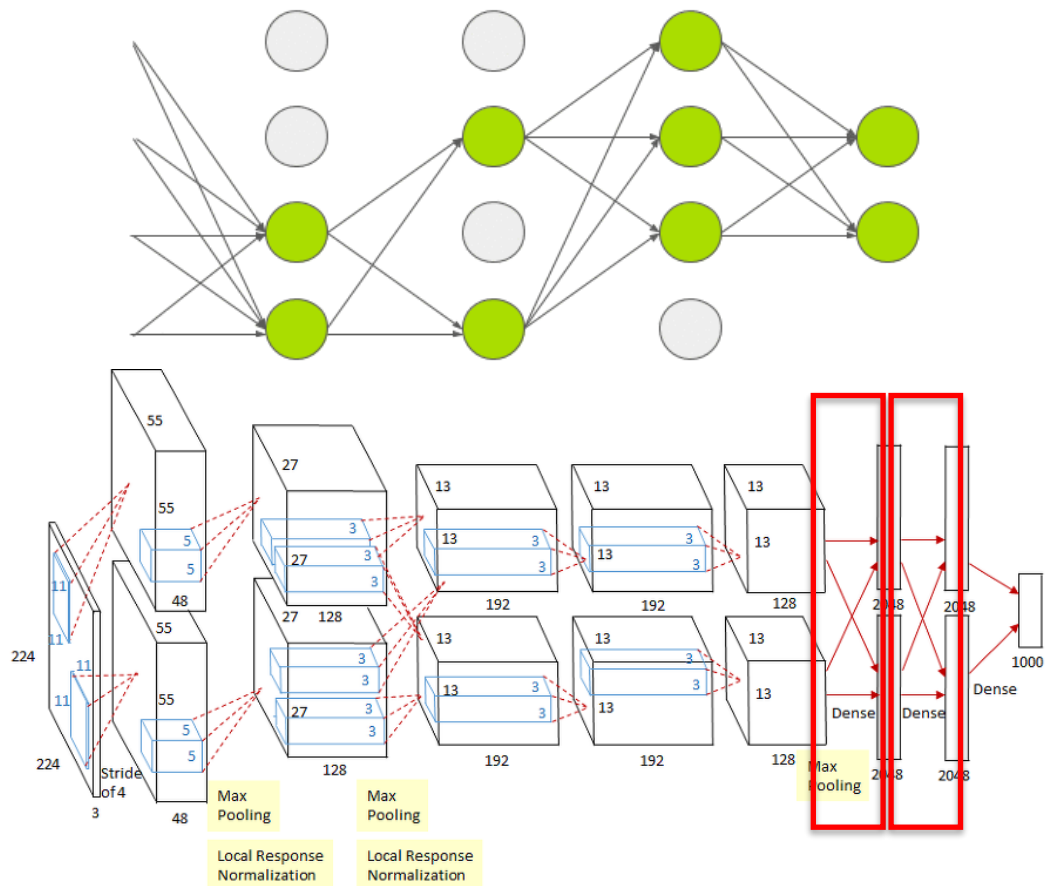
# Reducing Overfitting

- ▶ Data Augmentation by Random crops
  - ▶ AlexNet input images are of **size 227 by 227**, which are randomly sampled from ImageNet's 256 by 256
    - ▶ The paper mentions the network inputs to be 224, but that is a mistake




# Reducing Overfitting

## ► Dropout




# AlexNet Architecture

- ▶ First large scale convolutional neural network
- ▶ 5 Convolution layers
- ▶ 3 Fully connected layers




# Feature Extractor

## ► 1<sup>st</sup> Convolutional Layer





$1 \times 227 \times 227 \times 3(\text{RGB})$



$1 \times 55 \times 55 \times 96$


# Feature Extractor

## ► 2<sup>nd</sup> Convolutional Layer



# Feature Extractor


## ► 3<sup>rd</sup> Convolutional Layer






# Feature Extractor

## ► 4<sup>th</sup> Convolutional Layer




# Feature Extractor

## ► 5<sup>th</sup> Convolutional Layer






# Classifier



## ► 1<sup>st</sup> Fully Connected Layer



## ► 2<sup>nd</sup> Fully Connected Layer



## ► Output Layer



# AlexNet

## ► Model

```
model = keras.models.Sequential([
    keras.layers.Conv2D(filters=96, kernel_size=(11,11), strides=(4,4), activation='relu', input_shape=(227,227,3)),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPool2D(pool_size=(3,3), strides=(2,2)),
    keras.layers.Conv2D(filters=256, kernel_size=(5,5), strides=(1,1), activation='relu', padding="same"),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPool2D(pool_size=(3,3), strides=(2,2)),
    keras.layers.Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), activation='relu', padding="same"),
    keras.layers.BatchNormalization(),
    keras.layers.Conv2D(filters=384, kernel_size=(3,3), strides=(1,1), activation='relu', padding="same"),
    keras.layers.BatchNormalization(),
    keras.layers.Conv2D(filters=256, kernel_size=(3,3), strides=(1,1), activation='relu', padding="same"),
    keras.layers.BatchNormalization(),
    keras.layers.MaxPool2D(pool_size=(3,3), strides=(2,2)),
    keras.layers.Flatten(),
    keras.layers.Dense(4096, activation='relu'),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(4096, activation='relu'),
    keras.layers.Dropout(0.5),
    keras.layers.Dense(10, activation='softmax')
])
```

# Future work

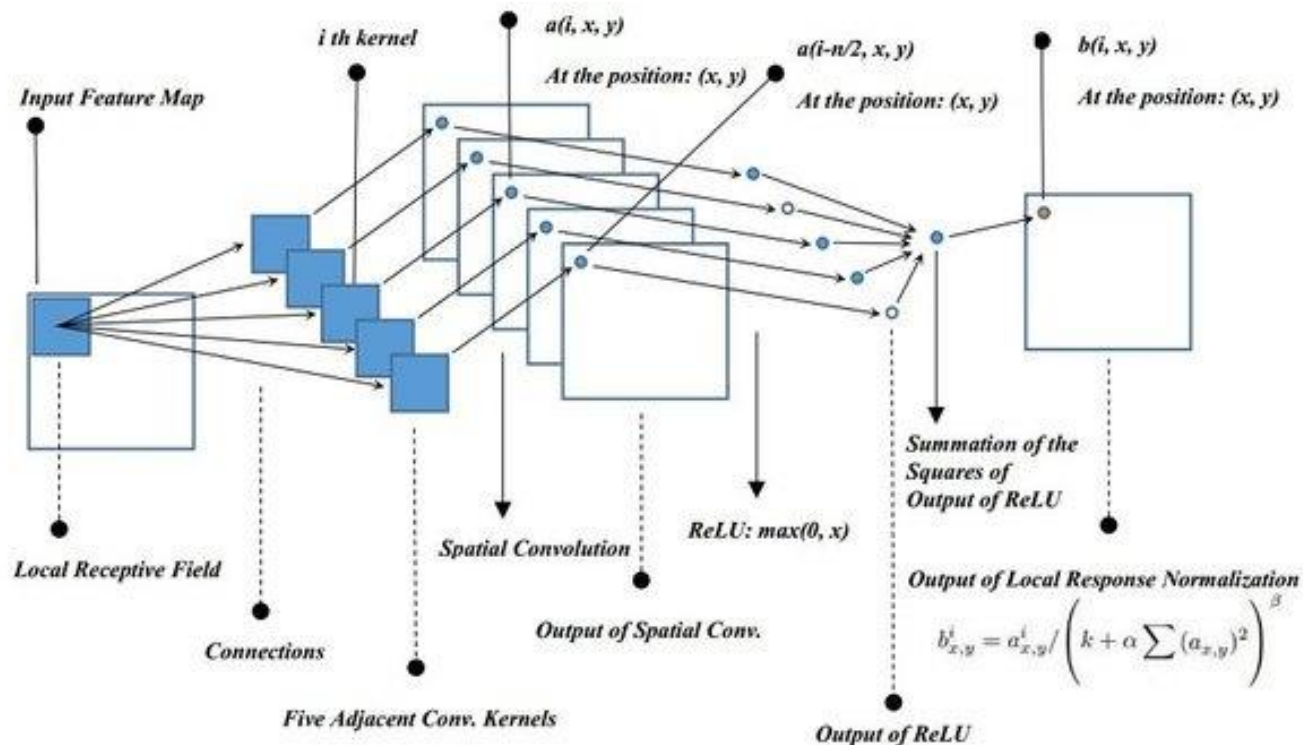
- ▶ Data Augmentation
  - ▶ Principal component Analysis (PCA)



# Future Work

## ► Local Response Normalization

□ ~: Batch Normalization



# Local Response Normalization

- ReLUs do not require input normalization to prevent them from saturating
- However, Local Response Normalization aids generalization

Activity of a neuron by applying  
kernel  $i$  at position  $(x,y)$

$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0, i-\frac{n}{2})}^{\min(N-1, i+\frac{n}{2})} (a_{x,y}^j)^2 \right)^\beta$$

$$\begin{aligned} k &= 2 \\ n &= 5 \\ \alpha &= 10^{-4} \\ \beta &= 0.75 \end{aligned}$$

- Improvement:
  - top-1 error rate by **1.4%**
  - top-5 error rate by **1.2%**

sum runs over  $n$  "adjacent" kernel  
maps at the same spatial position

# Architecture

## Local Response Normalization

- No need to input normalization with ReLUs.
- But still the following local normalization scheme helps generalization.

$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0,i-n/2)}^{\min(N-1,i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

Response-normalized activity

Activity of a neuron computed by applying kernel  $l$  at position  $(x,y)$  and then applying the ReLU nonlinearity

- Response normalization reduces top-1 and top-5 error rates by 1.4% and 1.2% , respectively.